

# Energy Load Forecasting

for the Global Energy Forecasting Competition 2014

Presentation of Semester Project  
by Fabian Brix



ÉCOLE POLYTECHNIQUE  
FÉDÉRALE DE LAUSANNE

January 27, 2015

- ▶ **Introduction**
- ▶ Data
- ▶ Features
- ▶ Forecasting Methods
- ▶ Quantile Forecasting
- ▶ Error Measures
- ▶ Method Evaluation
- ▶ Results

- ▶ Dates: 08/15/2014 to 12/15/2014
- ▶ Tracks: energy **load forecasting**, price forecasting, wind, solar
- ▶ Task: forecast the distribution, in quantiles, of the hourly district load of an energy utility.
- ▶ Includes: forecasting the temperature of different zones in the district.
- ▶ Forecasting horizon: 1 month
- ▶ Competition type: rolling forecast  
⇒ Target for preceding forecasting task published every week to prepare for next horizon

- ▶ Dates: 08/15/2014 to 12/15/2014
- ▶ Tracks: energy **load forecasting**, price forecasting, wind, solar
- ▶ Task: forecast the distribution, in quantiles, of the hourly district load of an energy utility.
- ▶ Includes: forecasting the temperature of different zones in the district.
- ▶ Forecasting horizon: 1 month
- ▶ Competition type: rolling forecast  
⇒ Target for preceding forecasting task published every week to prepare for next horizon

- ▶ Dates: 08/15/2014 to 12/15/2014
- ▶ Tracks: energy **load forecasting**, price forecasting, wind, solar
- ▶ Task: forecast the distribution, in quantiles, of the hourly district load of an energy utility.
- ▶ Includes: forecasting the temperature of different zones in the district.
- ▶ Forecasting horizon: 1 month
- ▶ Competition type: rolling forecast  
⇒ Target for preceding forecasting task published every week to prepare for next horizon

- ▶ Dates: 08/15/2014 to 12/15/2014
- ▶ Tracks: energy **load forecasting**, price forecasting, wind, solar
- ▶ Task: forecast the distribution, in quantiles, of the hourly district load of an energy utility.
- ▶ Includes: forecasting the temperature of different zones in the district.
- ▶ Forecasting horizon: 1 month
- ▶ Competition type: rolling forecast  
⇒ Target for preceding forecasting task published every week to prepare for next horizon

- ▶ Dates: 08/15/2014 to 12/15/2014
- ▶ Tracks: energy **load forecasting**, price forecasting, wind, solar
- ▶ Task: forecast the distribution, in quantiles, of the hourly district load of an energy utility.
- ▶ Includes: forecasting the temperature of different zones in the district.
- ▶ Forecasting horizon: 1 month
- ▶ Competition type: rolling forecast  
⇒ Target for preceding forecasting task published every week to prepare for next horizon


- ▶ Dates: 08/15/2014 to 12/15/2014
- ▶ Tracks: energy **load forecasting**, price forecasting, wind, solar
- ▶ Task: forecast the distribution, in quantiles, of the hourly district load of an energy utility.
- ▶ Includes: forecasting the temperature of different zones in the district.
- ▶ Forecasting horizon: 1 month
- ▶ Competition type: rolling forecast  
⇒ Target for preceding forecasting task published every week to prepare for next horizon



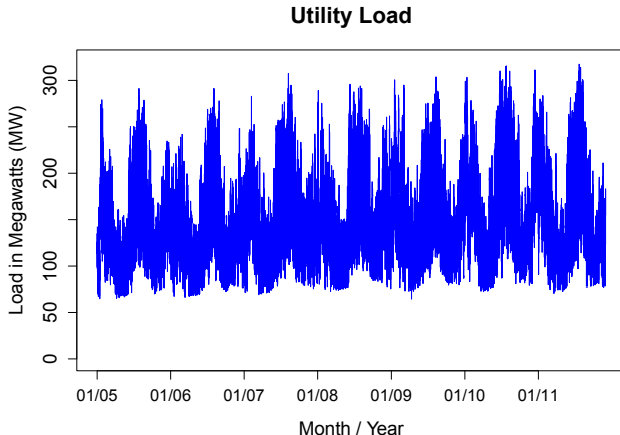
- ▶ Dates: 08/15/2014 to 12/15/2014
- ▶ Tracks: energy **load forecasting**, price forecasting, wind, solar
- ▶ Task: forecast the distribution, in quantiles, of the hourly district load of an energy utility.
- ▶ Includes: forecasting the temperature of different zones in the district.
- ▶ Forecasting horizon: 1 month
- ▶ Competition type: rolling forecast  
⇒ Target for preceding forecasting task published every week to prepare for next horizon

- ▶ Introduction
- ▶ **Data**
- ▶ Features
- ▶ Forecasting Methods
- ▶ Quantile Forecasting
- ▶ Error Measures
- ▶ Method Evaluation
- ▶ Results

- ▶ zonal: 25 series of hourly temperature (Fahrenheit)
- ▶ 01/01/2001 1am to 12/01/2011 midnight.



- ▶ district level: hourly load (Megawatts)
- ▶ 01/01/2005 1am to 12/01/2011 midnight.



- ▶ Introduction
- ▶ Data
- ▶ **Features**
- ▶ Forecasting Methods
- ▶ Quantile Forecasting
- ▶ Error Measures
- ▶ Method Evaluation
- ▶ Results

For temperature & load:

- ▶ YLAG:  $y_{-365}$  days
- ▶ DLAG:  $y_{-35}$  days
- ▶ TOY: time of year

For load only:

- ▶ CTEMP: current temperature
- ▶ DAYT: day type (holiday!)

Type	Mon	Tue	Wed	Thu	Fri	Sat	Sun	Holiday!
SDAYT	2	2	2	2	2	1	1	-
WDAYT	2	3	4	5	6	7	1	-
DAYT	2	3	4	5	6	7	1	8

For temperature & load:

- ▶ YLAG:  $y_{-365}$  days
- ▶ DLAG:  $y_{-35}$  days
- ▶ TOY: time of year

For load only:

- ▶ CTEMP: current temperature
- ▶ DAYT: day type (holiday!)

Type	Mon	Tue	Wed	Thu	Fri	Sat	Sun	Holiday!
SDAYT	2	2	2	2	2	1	1	-
WDAYT	2	3	4	5	6	7	1	-
DAYT	2	3	4	5	6	7	1	8

For temperature & load:

- ▶ YLAG:  $y_{-365}$  days
- ▶ DLAG:  $y_{-35}$  days
- ▶ TOY: time of year

For load only:

- ▶ CTEMP: current temperature
- ▶ DAYT: day type (holiday!)

Type	Mon	Tue	Wed	Thu	Fri	Sat	Sun	Holiday!
SDAYT	2	2	2	2	2	1	1	-
WDAYT	2	3	4	5	6	7	1	-
DAYT	2	3	4	5	6	7	1	8



- ▶ Introduction
- ▶ Data
- ▶ Features
- ▶ **Forecasting Methods**
- ▶ Quantile Forecasting
- ▶ Error Measures
- ▶ Method Evaluation
- ▶ Results

- ▶ Additive predictor with non-linear functions
- ▶  $y = s_1(x_1) + s_2(x_2) + \dots + s_p(x_p) + \varepsilon$
- ▶  $s_i(x_i)$  are smooth functions estimated from the data  $x_i$
- ▶ Noise  $\varepsilon$  assumed to be Gaussian
- ▶ Use of standard configuration of GAM in R package “mgcv”

- ▶ Additive predictor with non-linear functions
- ▶  $y = s_1(x_1) + s_2(x_2) + \dots + s_p(x_p) + \varepsilon$
- ▶  $s_i(x_i)$  are smooth functions estimated from the data  $x_i$
- ▶ Noise  $\varepsilon$  assumed to be Gaussian
- ▶ Use of standard configuration of GAM in R package “mgcv”

- ▶ Additive predictor with non-linear functions
- ▶  $y = s_1(x_1) + s_2(x_2) + \dots + s_p(x_p) + \varepsilon$
- ▶  $s_i(x_i)$  are smooth functions estimated from the data  $x_i$
- ▶ Noise  $\varepsilon$  assumed to be Gaussian
- ▶ Use of standard configuration of GAM in R package “mgcv”

- ▶ Additive predictor with non-linear functions
- ▶  $y = s_1(x_1) + s_2(x_2) + \dots + s_p(x_p) + \varepsilon$
- ▶  $s_i(x_i)$  are smooth functions estimated from the data  $x_i$
- ▶ Noise  $\varepsilon$  assumed to be Gaussian
- ▶ Use of standard configuration of GAM in R package “mgcv”

- ▶ Additive predictor with non-linear functions
- ▶  $y = s_1(x_1) + s_2(x_2) + \dots + s_p(x_p) + \varepsilon$
- ▶  $s_i(x_i)$  are smooth functions estimated from the data  $x_i$
- ▶ Noise  $\varepsilon$  assumed to be Gaussian
- ▶ Use of standard configuration of GAM in R package “mgcv”

## Configuration using R package nnet

- ▶ Hidden layer: 1
- ▶ Hidden units: 5, 10, 15, 20, 25, 30
- ▶ Linear output for regression

## Configuration using R package nnet

- ▶ Hidden layer: 1
- ▶ Hidden units: 5, 10, 15, 20, 25, 30
- ▶ Linear output for regression



Configuration using R package nnet


- ▶ Hidden layer: 1
- ▶ Hidden units: 5, 10, 15, 20, 25, 30
- ▶ Linear output for regression

Configuration using R package nnet

- ▶ Hidden layer: 1
- ▶ Hidden units: 5, 10, 15, 20, 25, 30
- ▶ Linear output for regression

Configuration using R package nnet

- ▶ Hidden layer: 1
- ▶ Hidden units: 5, 10, 15, 20, 25, 30
- ▶ Linear output for regression



## Configuration using R package randomForest


- ▶ # trees: 50, 100, 150, 300

Configuration using R package randomForest

- ▶ # trees: 50, 100, 150, 300

Configuration using R package randomForest

- ▶ # trees: 50, 100, 150, 300



- ▶ Model: hyperparameter settings + formula
- ▶ Formula: collection of features used in model (R)  
⇒ Create point forecasts with different forecasting methods

- ▶ Model: hyperparameter settings + formula
- ▶ Formula: collection of features used in model (R)
  - ⇒ Create point forecasts with different forecasting methods



- ▶ Model: hyperparameter settings + formula
- ▶ Formula: collection of features used in model (R)
  - ⇒ Create point forecasts with different forecasting methods

- ▶ Introduction
- ▶ Data
- ▶ Features
- ▶ Forecasting Methods
- ▶ **Quantile Forecasting**
- ▶ Error Measures
- ▶ Method Evaluation
- ▶ Results

- ▶ Future: random variable with “forecast distribution”
- ▶ “Quantile forecast”: quantiles of this distribution (percentiles)
- ▶ Point forecasts correspond to mean of forecast distribution  
→ Offset **quantiles of training residuals** by point forecast


- ▶ Future: random variable with “forecast distribution”
- ▶ “Quantile forecast”: quantiles of this distribution (percentiles)
- ▶ Point forecasts correspond to mean of forecast distribution  
→ Offset **quantiles of training residuals** by point forecast

- ▶ Future: random variable with “forecast distribution”
- ▶ “Quantile forecast”: quantiles of this distribution (percentiles)
- ▶ Point forecasts correspond to mean of forecast distribution
  - Offset **quantiles of training residuals** by point forecast

- ▶ Future: random variable with “forecast distribution”
- ▶ “Quantile forecast”: quantiles of this distribution (percentiles)
- ▶ Point forecasts correspond to mean of forecast distribution  
→ Offset **quantiles of training residuals** by point forecast

- ▶ Future: random variable with “forecast distribution”
- ▶ “Quantile forecast”: quantiles of this distribution (percentiles)
- ▶ Point forecasts correspond to mean of forecast distribution  
→ Offset **quantiles of training residuals** by point forecast

- ▶ Future: random variable with “forecast distribution”
- ▶ “Quantile forecast”: quantiles of this distribution (percentiles)
- ▶ Point forecasts correspond to mean of forecast distribution  
→ Offset **quantiles of training residuals** by point forecast





- ▶ Introduction
- ▶ Data
- ▶ Features
- ▶ Forecasting Methods
- ▶ Quantile Forecasting
- ▶ **Error Measures**
- ▶ Method Evaluation
- ▶ Results

- ▶ Used to evaluate goodness of fit
- ▶ Scale-independent (%)

$$\text{MAPE} = \frac{100}{N} \sum_{i=1}^N \left| \frac{\xi_i}{y_i} \right|$$

- ▶ residual:  $\xi_i = y_i - \hat{y}_i$
- ▶  $y_i$ : time series value at time  $i$
- ▶  $\hat{y}_i$ : forecast at time  $i$

- ▶ Used to evaluate goodness of fit
- ▶ Scale-independent (%)

$$\text{MAPE} = \frac{100}{N} \sum_{i=1}^N \left| \frac{\xi_i}{y_i} \right|$$

- ▶ residual:  $\xi_i = y_i - \hat{y}_i$
- ▶  $y_i$ : time series value at time  $i$
- ▶  $\hat{y}_i$ : forecast at time  $i$

- ▶ Used to evaluate goodness of fit
- ▶ Scale-independent (%)

$$\text{MAPE} = \frac{100}{N} \sum_{i=1}^N \left| \frac{\xi_i}{y_i} \right|$$

- ▶ residual:  $\xi_i = y_i - \hat{y}_i$
- ▶  $y_i$ : time series value at time  $i$
- ▶  $\hat{y}_i$ : forecast at time  $i$

- ▶ Used to evaluate goodness of fit
- ▶ Scale-independent (%)

$$\text{MAPE} = \frac{100}{N} \sum_{i=1}^N \left| \frac{\xi_i}{y_i} \right|$$

- ▶ residual:  $\xi_i = y_i - \hat{y}_i$
- ▶  $y_i$ : time series value at time  $i$
- ▶  $\hat{y}_i$ : forecast at time  $i$

- ▶ Used to evaluate goodness of fit
- ▶ Scale-independent (%)

$$\text{MAPE} = \frac{100}{N} \sum_{i=1}^N \left| \frac{\xi_i}{y_i} \right|$$

- ▶ residual:  $\xi_i = y_i - \hat{y}_i$
- ▶  $y_i$ : time series value at time  $i$
- ▶  $\hat{y}_i$ : forecast at time  $i$

- ▶ Used to evaluate quantile forecasts
- ▶ And measure performance against competition leaderboard
- ▶ Strictly positive, the lower the more accurate the forecast
- ▶ Score at  $i$  given by  $\text{mean}(L_{0.01}, L_{0.02}, \dots, L_{0.99})$

$$L_{\tau}(\xi) = \begin{cases} \tau\xi & \text{if } \xi \geq 0 \\ (\tau - 1)\xi & \text{if } \xi < 0 \end{cases} \quad \text{where } \xi = (y_i - \hat{y}_{i,\tau})$$

- ▶ Used to evaluate quantile forecasts
- ▶ And measure performance against competition leaderboard
- ▶ Strictly positive, the lower the more accurate the forecast
- ▶ Score at  $i$  given by  $\text{mean}(L_{0.01}, L_{0.02}, \dots, L_{0.99})$

$$L_{\tau}(\xi) = \begin{cases} \tau\xi & \text{if } \xi \geq 0 \\ (\tau - 1)\xi & \text{if } \xi < 0 \end{cases} \quad \text{where } \xi = (y_i - \hat{y}_{i,\tau})$$




- ▶ Used to evaluate quantile forecasts
- ▶ And measure performance against competition leaderboard
- ▶ Strictly positive, the lower the more accurate the forecast
- ▶ Score at  $i$  given by  $\text{mean}(L_{0.01}, L_{0.02}, \dots, L_{0.99})$

$$L_{\tau}(\xi) = \begin{cases} \tau\xi & \text{if } \xi \geq 0 \\ (\tau - 1)\xi & \text{if } \xi < 0 \end{cases} \quad \text{where } \xi = (y_i - \hat{y}_{i,\tau})$$

- ▶ Used to evaluate quantile forecasts
- ▶ And measure performance against competition leaderboard
- ▶ Strictly positive, the lower the more accurate the forecast
- ▶ Score at  $i$  given by  $\text{mean}(L_{0.01}, L_{0.02}, \dots, L_{0.99})$

$$L_{\tau}(\xi) = \begin{cases} \tau\xi & \text{if } \xi \geq 0 \\ (\tau - 1)\xi & \text{if } \xi < 0 \end{cases} \quad \text{where } \xi = (y_i - \hat{y}_{i,\tau})$$



- ▶ Introduction
- ▶ Data
- ▶ Features
- ▶ Forecasting Methods
- ▶ Quantile Forecasting
- ▶ Error Measures
- ▶ **Method Evaluation**
- ▶ Results

- ▶ Use whole competition dataset
- ▶ Initial data + 14 months from subsequent tasks
- ▶ Use rolling forecast of 14 month period to evaluate forecasts

- ▶ Use whole competition dataset
- ▶ Initial data + 14 months from subsequent tasks
- ▶ Use rolling forecast of 14 month period to evaluate forecasts

- ▶ Use whole competition dataset
- ▶ Initial data + 14 months from subsequent tasks
- ▶ Use rolling forecast of 14 month period to evaluate forecasts



- ▶ Use whole competition dataset
- ▶ Initial data + 14 months from subsequent tasks
- ▶ Use rolling forecast of 14 month period to evaluate forecasts

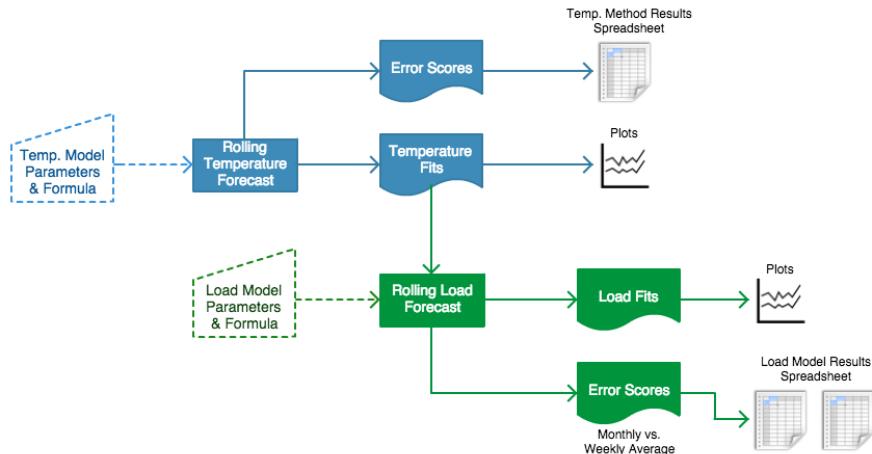


- ▶ Use whole competition dataset
- ▶ Initial data + 14 months from subsequent tasks
- ▶ Use rolling forecast of 14 month period to evaluate forecasts





# Workflow Diagram




- ▶ Introduction
- ▶ Data
- ▶ Features
- ▶ Forecasting Methods
- ▶ Quantile Forecasting
- ▶ Error Measures
- ▶ Method Evaluation
- ▶ **Results**

- ▶ Best score: 12.04% MAPE
- ▶ Model: neural network and 30 hidden units

- ▶ Best score: 12.04% MAPE
- ▶ Model: neural network and 30 hidden units



# Results - Temperature

- ▶ Best score: 12.04% MAPE
- ▶ Model: neural network and 30 hidden units



# Results - Temperature

- ▶ Best score: 12.04% MAPE
- ▶ Model: neural network and 30 hidden units



- ▶ Effect of temperature  
→ Significant: -2.5% MAPE
- ▶ No effect of different temperature preprocessing (average, PCA, 1 series)
- ▶ Predicting temperature for training period of load forecast  
→ Pinball error improves
- ▶ Effect of forecast horizon  
→ results improved with 1 week horizon
- ▶ Effect of daytypes  
→ better results without

- ▶ Effect of temperature  
→ Significant: -2.5% MAPE
- ▶ No effect of different temperature preprocessing (average, PCA, 1 series)
- ▶ Predicting temperature for training period of load forecast  
→ Pinball error improves
- ▶ Effect of forecast horizon  
→ results improved with 1 week horizon
- ▶ Effect of daytypes  
→ better results without



- ▶ Effect of temperature
  - Significant: -2.5% MAPE
- ▶ No effect of different temperature preprocessing (average, PCA, 1 series)
- ▶ Predicting temperature for training period of load forecast
  - Pinball error improves
- ▶ Effect of forecast horizon
  - results improved with 1 week horizon
- ▶ Effect of daytypes
  - better results without


- ▶ Effect of temperature  
→ Significant: -2.5% MAPE
- ▶ No effect of different temperature preprocessing (average, PCA, 1 series)
- ▶ Predicting temperature for training period of load forecast  
→ Pinball error improves
- ▶ Effect of forecast horizon  
→ results improved with 1 week horizon
- ▶ Effect of daytypes  
→ better results without

- ▶ Effect of temperature  
→ Significant: -2.5% MAPE
- ▶ No effect of different temperature preprocessing (average, PCA, 1 series)
- ▶ Predicting temperature for training period of load forecast  
→ Pinball error improves
- ▶ Effect of forecast horizon  
→ results improved with 1 week horizon
- ▶ Effect of daytypes  
→ better results without

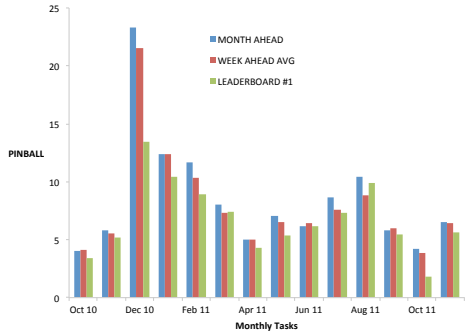

- ▶ Random forest, 100 trees, 1 week f. horizon
- ▶ MAPE: 16.1%, Pinball: 7.999

- ▶ Random forest, 100 trees, 1 week f. horizon
- ▶ MAPE: 16.1%, Pinball: 7.999

- ▶ Random forest, 100 trees, 1 week f. horizon
- ▶ MAPE: 16.1%, Pinball: 7.999



- ▶ Random forest, 100 trees, 1 week f. horizon
- ▶ MAPE: 16.1%, Pinball: 7.999




# Best Configuration - Final Leaderboard



ÉCOLE POLYTECHNIQUE  
FÉDÉRALE DE LAUSANNE

- ▶ Trimmed pinball mean: 7.16  $\Rightarrow$  #3
- ▶ Average position on task leaderboards: 5.64
- ▶ Not every contestant reaches equally high scores every week





- ▶ Find a development environment which makes debugging easier (R in vim)
- ▶ Analyze training residuals → white noise
- ▶ Further research on random forests for regression
- ▶ Produce and evaluate ensemble forecasts

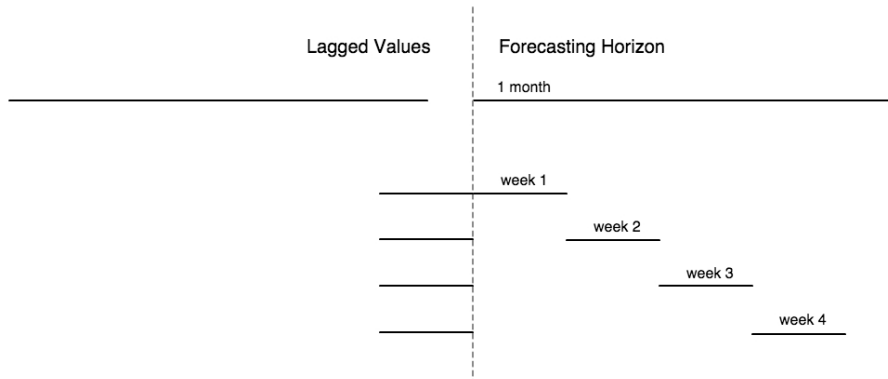
- ▶ Find a development environment which makes debugging easier (R in vim)
- ▶ Analyze training residuals → white noise
- ▶ Further research on random forests for regression
- ▶ Produce and evaluate ensemble forecasts



- ▶ Find a development environment which makes debugging easier (R in vim)
- ▶ Analyze training residuals → white noise
- ▶ Further research on random forests for regression
- ▶ Produce and evaluate ensemble forecasts

- ▶ Find a development environment which makes debugging easier (R in vim)
- ▶ Analyze training residuals → white noise
- ▶ Further research on random forests for regression
- ▶ Produce and evaluate ensemble forecasts



- ▶ Find a development environment which makes debugging easier (R in vim)
- ▶ Analyze training residuals → white noise
- ▶ Further research on random forests for regression
- ▶ Produce and evaluate ensemble forecasts

- ▶ Forecasts with two different time horizons: 1 month vs 1 week
- ▶ Forecasts of week 1 to 4 are combined + rest of monthly forecast





(a) time lag: 7 days, 168 hours (b) time lag: 7 days, 168 hours



(c) lag: 7 days, 168 hours

(d) lag: 35 days, 840 hours

